

1. Which of the following do you typically see in a ConvNet? (Check all that apply.)

1 / 1 point

✓ Correct

2. In order to be able to build very deep networks, we usually only use pooling layers to downsize the height/width of the activation volumes while convolutions are used with “valid” padding. Otherwise, we would downsize the input of the model too quickly.

1 / 1 point

✓ Correct

3. Training a deeper network (for example, adding additional layers to the network) allows the network to fit more complex functions and thus almost always results in lower training error. For this question, assume we’re referring to “plain” networks.

1 / 1 point

✓ Correct

4. The following equation captures the computation in a ResNet block. What goes into the two blanks above?

1 / 1 point

$$a^{[l+2]} = g(W^{[l+2]} g(W^{[l+1]} a^{[l]} + b^{[l+1]}) + b^{l+2} + \text{_____} ) + \text{_____}$$

✓ Correct

5. Which ones of the following statements on Residual Networks are true? (Check all that apply.)

1 / 1 point

✓ Correct

6. Suppose you have an input volume of dimension  $n_H \times n_W \times n_C$ . Which of the following statements you agree with? (Assume that “1x1 convolutional layer” below always uses a stride of 1 and no padding.)

1 / 1 point

✓ Correct

7. Which ones of the following statements on Inception Networks are true? (Check all that apply.)

1 / 1 point

✓ Correct

8. Which of the following are common reasons for using open-source implementations of ConvNets (both the model and/or weights)? Check all that apply.

1 / 1 point

✓ Correct

9. In Depthwise Separable Convolution you:

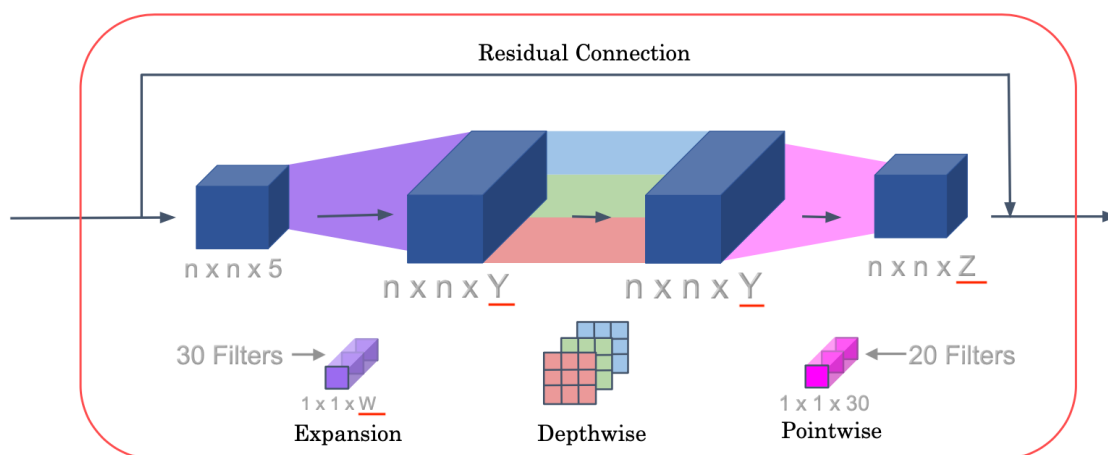
1 / 1 point

✓ Correct

10. Fill in the missing dimensions shown in the image below (marked  $W$ ,  $Y$ ,  $Z$ ).

1 / 1 point

## MobileNet v2 Bottleneck



✓ Correct